


Operating Systems

COMS(3010A)

Paging





Branden Ingram

branden.ingram@wits.ac.za

Recap

- **Memory Virtualisation**
- **Address Translation**
- **Segmentation**
- **Free Space Management**






Paging

- In Operating Systems, Paging is a storage mechanism used to retrieve processes from the secondary storage into the main memory in the form of pages.



Paging

- In Operating Systems, Paging is a storage mechanism used to retrieve processes from the secondary storage into the main memory in the form of pages.
- The main idea behind the paging is to divide each process(**Address Space**) into a fixed-sized unit called a **page**



Paging

- In Operating Systems, Paging is a storage mechanism used to retrieve processes from the secondary storage into the main memory in the form of pages.
- The main idea behind the paging is to divide each process(**Address Space**) into a fixed-sized unit called a **page**
 - Segmentation: variable size of logical segments(code, stack, heap, etc.)

Paging

- In Operating Systems, Paging is a storage mechanism used to retrieve processes from the secondary storage into the main memory in the form of pages.
- The main idea behind the paging is to **split** each process(**Address Space**) into a fixed-sized unit called a **page**
 - Segmentation: variable size of logical segments(code, stack, heap, etc.)
- With paging, **physical memory** is also **split** into some number of pages called a **page frame**.

Paging Translation ?

- **Page table** per process is needed **to translate** the virtual address to physical address.
- Similar to the segment register

Segment Register

| <u>Segment</u> | <u>Base</u> | <u>Size</u> |
|----------------|-------------|-------------|
| Code | 256KB | 4KB |
| Heap | 260KB | 60KB |
| Stack | 128KB | 64KB |

Paging Translation ?

- **Page table** per process is needed **to translate** the virtual address to physical address.
- Similar to the segment register

Segment Register


| Segment | Base | Size |
|---------|-------|------|
| Code | 256KB | 4KB |
| Heap | 260KB | 60KB |
| Stack | 128KB | 64KB |

Page Table

| Page | Page Frame |
|------|------------|
| 0 | 3 |
| 1 | 7 |
| 2 | 5 |
| 3 | 2 |




Advantages of Paging

- **Flexibility:** Supporting the abstraction of address space effectively
 - Don't need assumption how heap and stack grow and are used.




Advantages of Paging

- **Flexibility:** Supporting the abstraction of address space effectively
 - Don't need assumption how heap and stack grow and are used.
- **Simplicity:** ease of free-space management
 - The page in address space and the page frame are the same size.
 - Easy to allocate and keep a free list




Paging : Example

- 128-byte physical memory with 16 bytes page frames
- 64-byte address space with 16 bytes pages



64 byte Address Space




**64 byte Address Space placed in
Physical Memory**

Address Translation : Example

- **Two components in the virtual address**


- **VPN**: virtual page number
- **Offset**: offset within the page




Address Translation : Example

- **Two components in the virtual address**

- **VPN**: virtual page number
- **Offset**: offset within the page

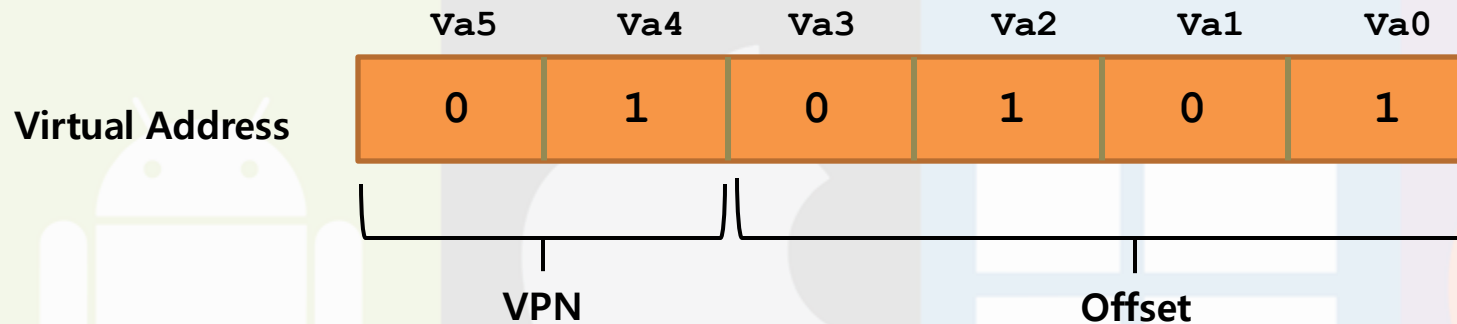


- **Example: virtual address 21 in 64-byte address space**



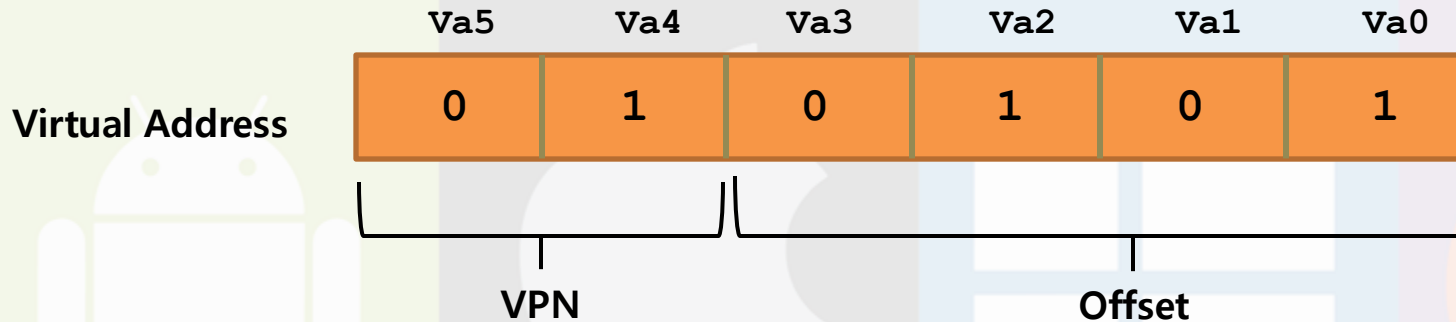
Address Translation : Example

- Example: virtual address 21 in 64-byte address space



Address Translation : Example

- Example: virtual address 21 in 64-byte address space




How do I know the Virtual Address is x bits ?

Address space= 64bytes

$\log_2(64) = 6$

Address Translation : Example

- Example: virtual address 21 in 64-byte address space



How do I know the Virtual Address is x bits ?

Address space = 64bytes

$\log_2(64) = 6$ (Virtual Address)

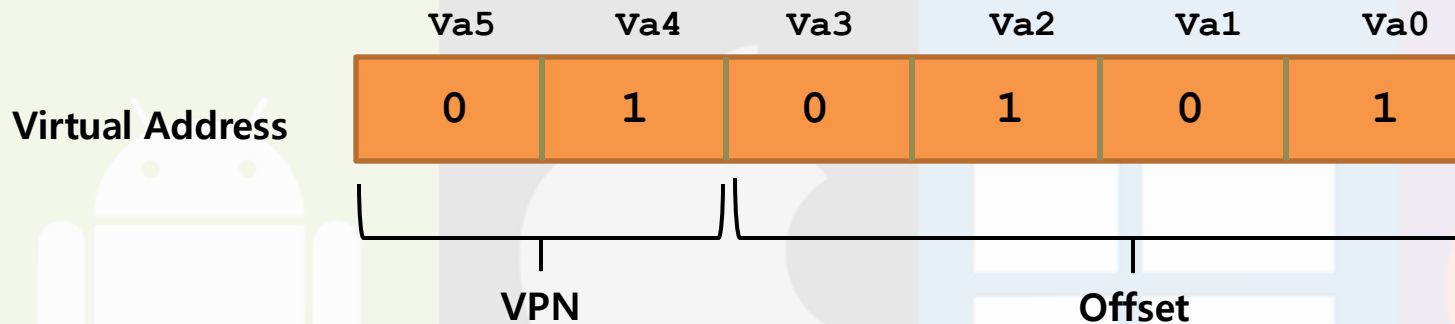
How do I know its 2 bits for VPN and 4 bits for Offset

Page Size = 16

$\log_2(16) = 4$ (offset)

Address Translation : Example

- Example: virtual address 21 in 64-byte address space



How do I know the Virtual Address is 6 bits ?

Address space = 64bytes

$\log_2(64) = 6$ (Virtual Address)

How do I know its 2 bits for VPN and 4 bits for Offset

Page Size = 16bytes

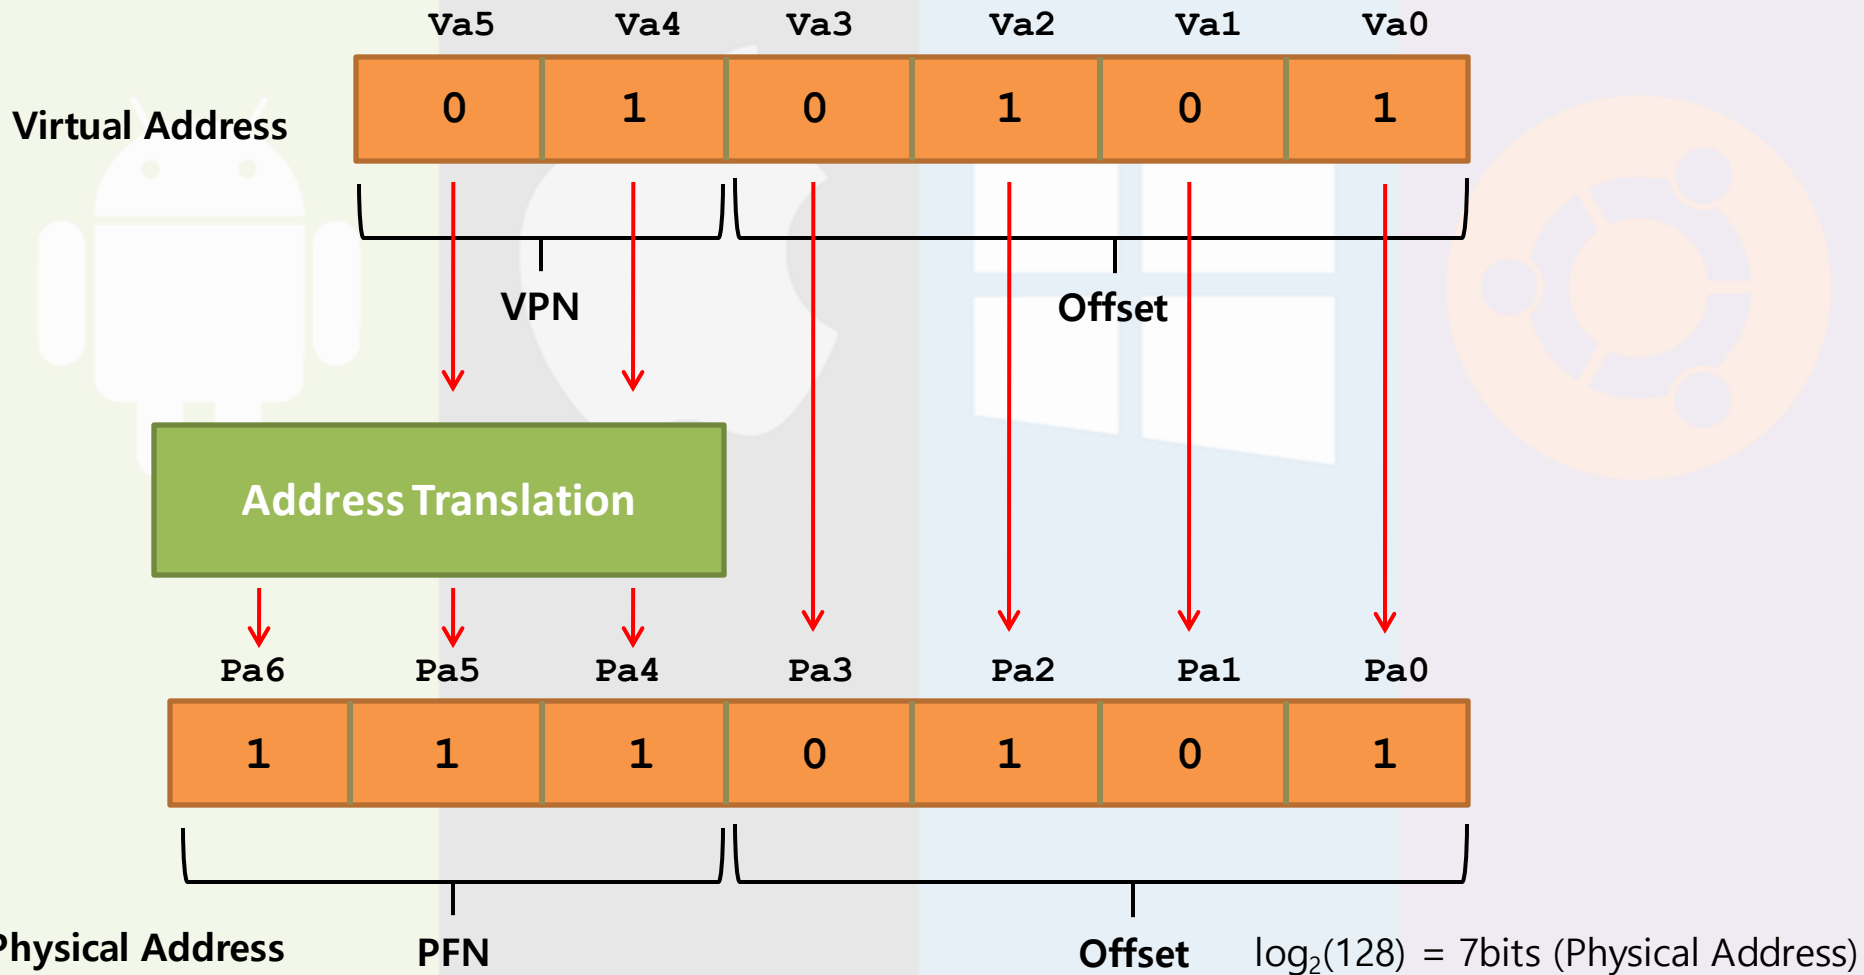
$\log_2(16) = 4$ (offset)

How do I know 21 = 010101

$21_{10} = 10101_2(5\text{bits}) = 010101_2(6\text{bits})$


Address Translation : Example

- Example: virtual address 21 in 64-byte address space



Paging : Example

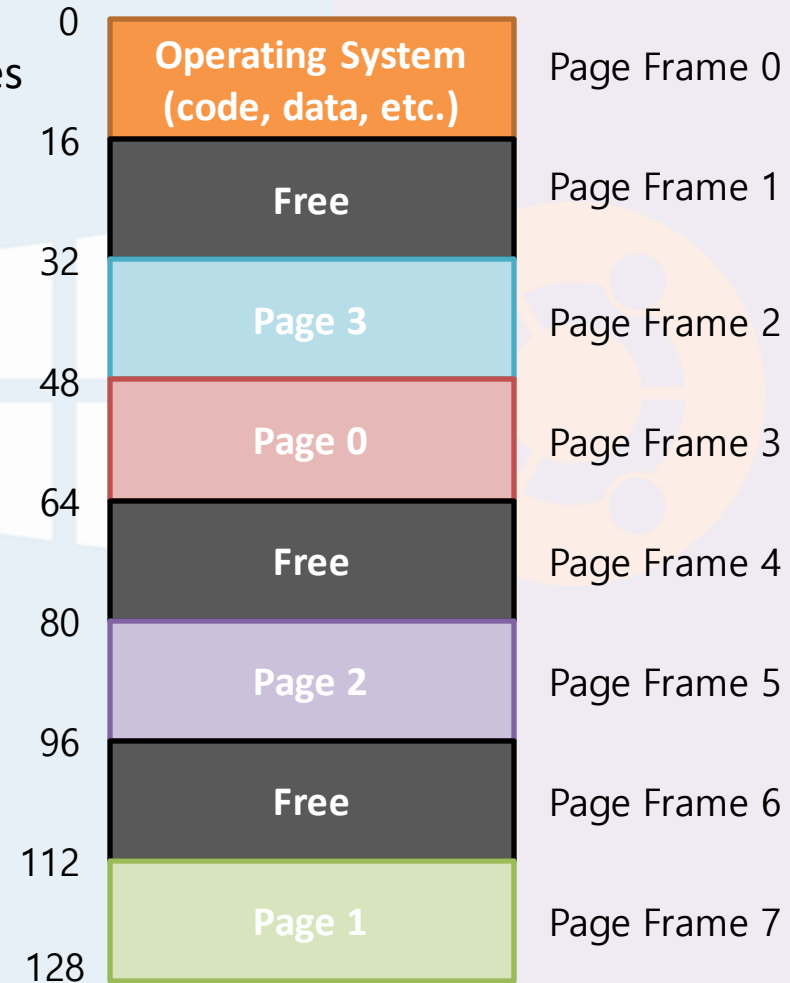
- 128-byte physical memory with 16 bytes page frames
- 64-byte address space with 16 bytes pages



64 byte Address Space

VPN = 01

PFN = 3bits



64 byte Address Space placed in Physical Memory

Paging : Example

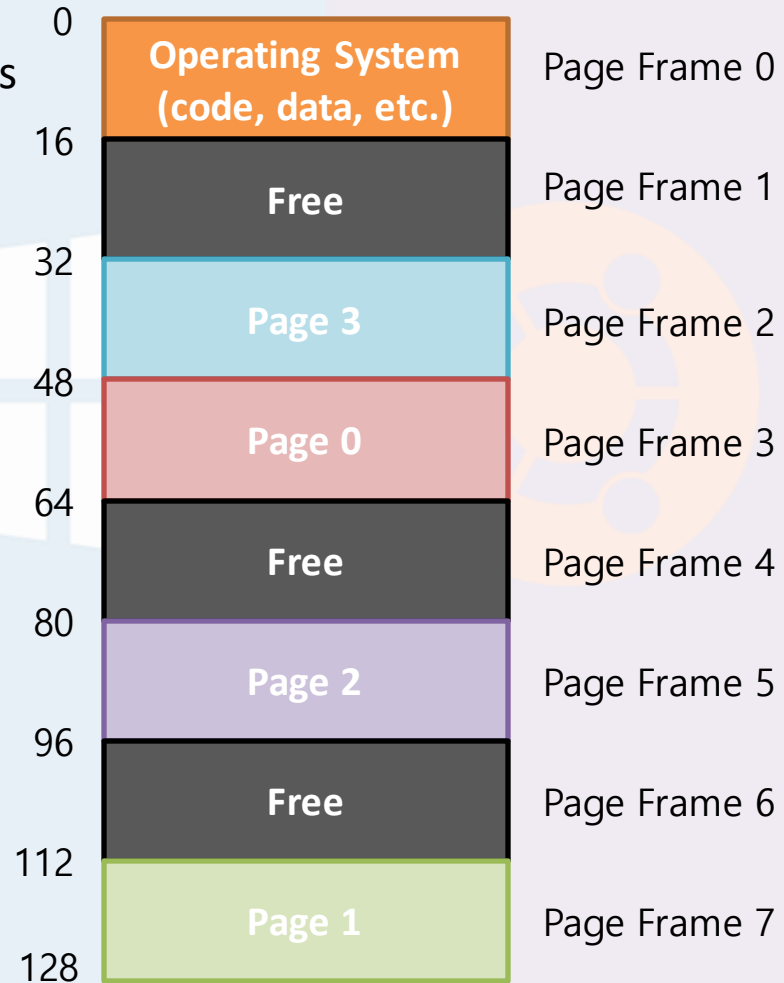
- 128-byte physical memory with 16 bytes page frames
- 64-byte address space with 16 bytes pages



64 byte Address Space

VPN = 01

PFN = 3bits



**64 byte Address Space placed in
Physical Memory**

Paging : Example

- 128-byte physical memory with 16 bytes page frames
- 64-byte address space with 16 bytes pages



64 byte Address Space

VPN = 01

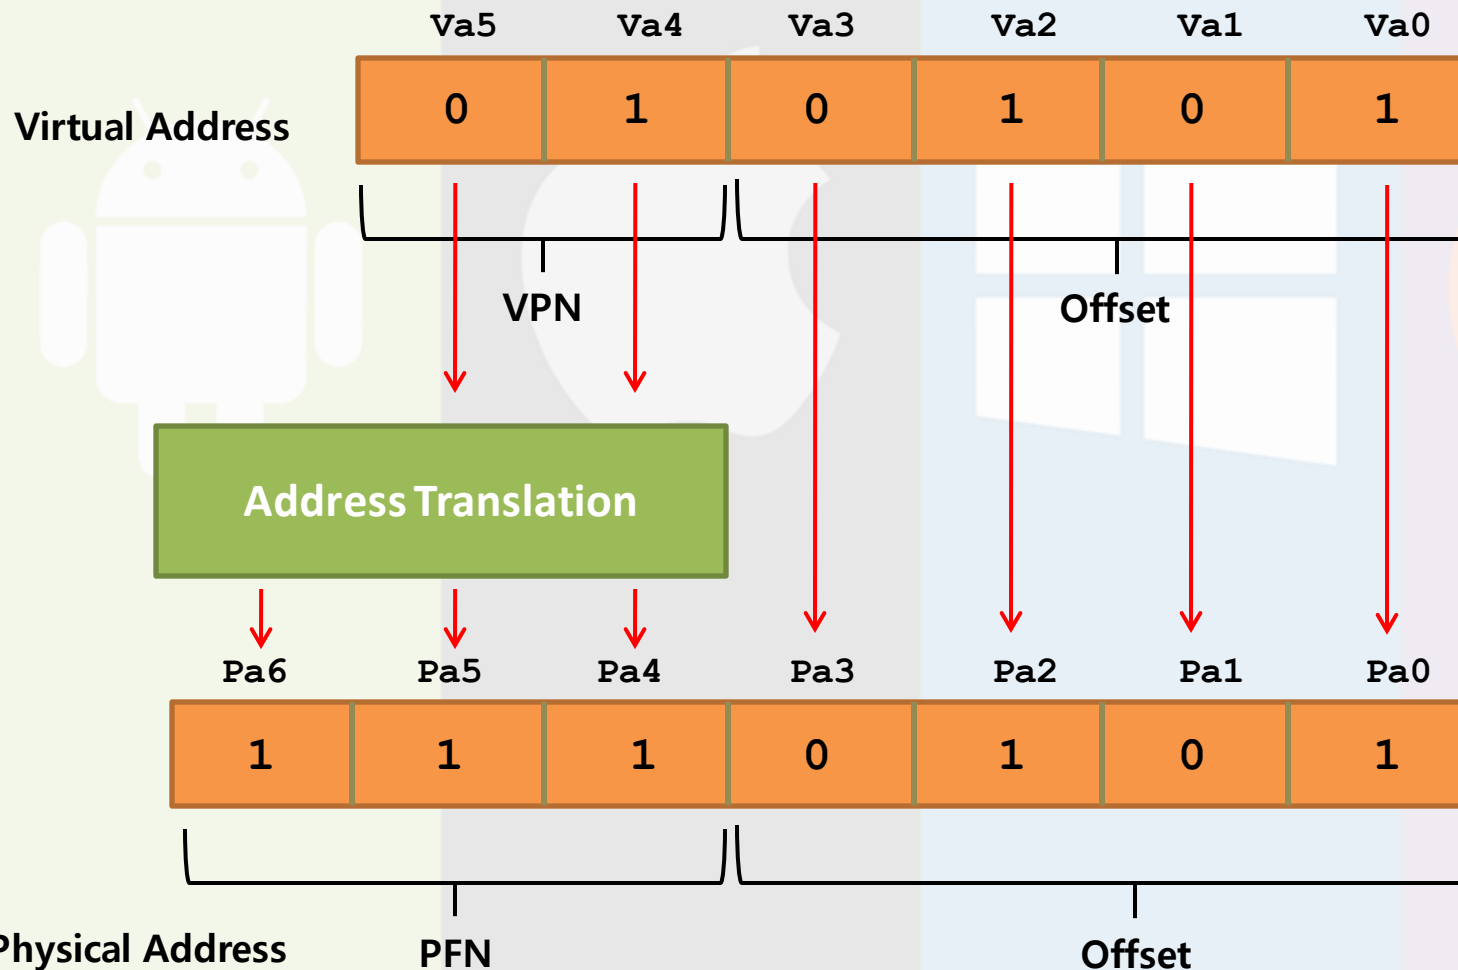
PFN = 3bits



64 byte Address Space placed in
Physical Memory

Address Translation : Example

- Example: virtual address 21 in 64-byte address space



Paging Translation ?




- **Page table** per process is needed **to translate** the virtual address to physical address.
- Similar to the segment register

Page Table

| Page | Page Frame |
|------|------------|
| 0 | 3 |
| 1 | 7 |
| 2 | 5 |
| 3 | 2 |

Ok well where are these tables stored?


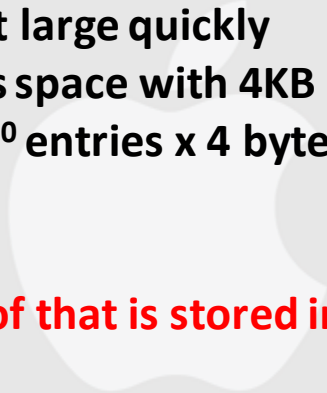

- Page tables can get large quickly
 - 32 bit address space with 4KB pages, 20bits for the VPN
 - $4\text{MB} = 2^{20}$ entries x 4 bytes per page table entry
- This is too large to store entries in dedicated CPU register



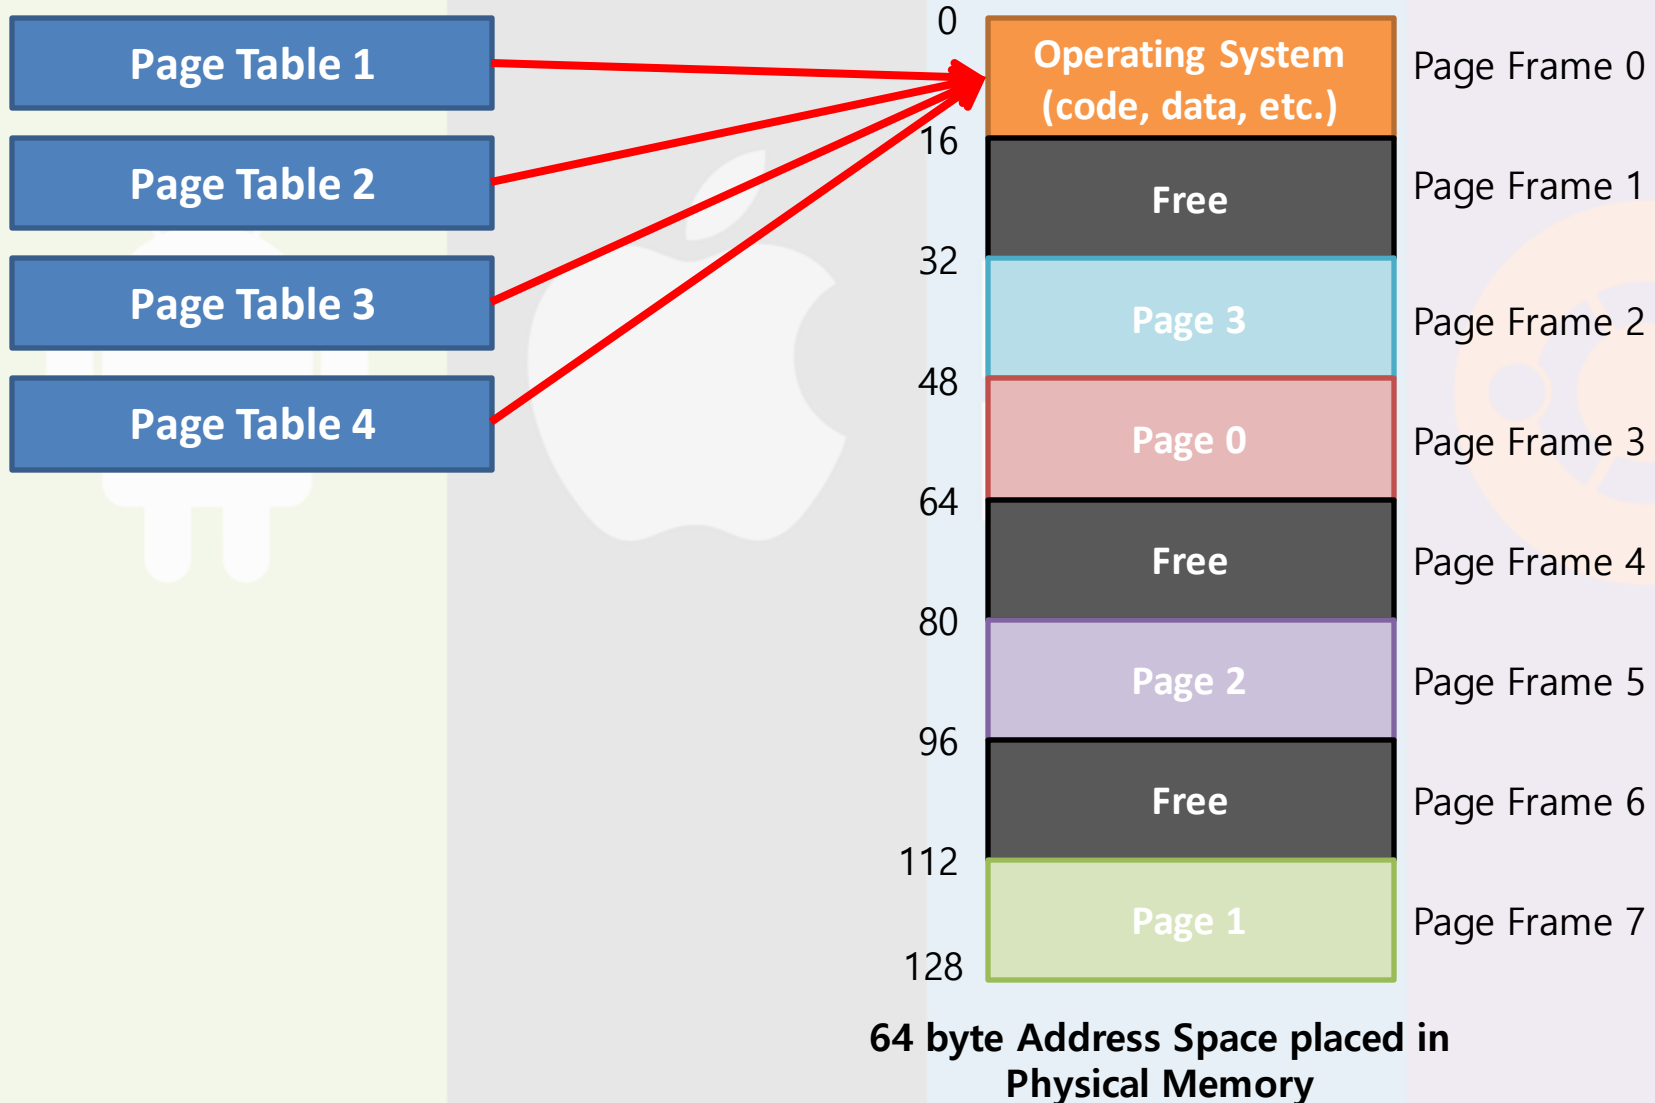
Ok well where are these tables stored?

- Page tables can get large quickly
 - 32 bit address space with 4KB pages, 20bits for the VPN
 - 4MB = 2^{20} entries x 4 bytes per page table entry

• And Therefore all of that is stored in memory



Page Table in Kernel Memory



What is in the Page Table?

- The page table is just a data structure that is used to map the virtual address to the physical address

Page Table

| Page | Page Frame |
|------|------------|
| 0 | 3 |
| 1 | 7 |
| 2 | 5 |
| 3 | 2 |

- In its simplest form it is just an array

What is in the Page Table?

- The page table is just a data structure that is used to map the virtual address to the physical address



Page Table

| Page | Page Frame |
|------|------------|
| 0 | 3 |
| 1 | 7 |
| 2 | 5 |
| 3 | 2 |

Page table Entry

- In its simplest form it is just an array
- The OS indexes the array by VPN and looks up the page table entry

Page Table Entry : Example x86



- P: present bit
- R/W: read/write bit
- U/S: supervisor
- A: accessed bit
- D: dirty bit
- PFN: the page frame number

Page Table Entry : Example x86



- **Valid Bit(P)**: Indicating whether the particular translation is valid.
- **Protection Bit(R/W)**: Indicating whether the page could be read from, written to, or executed from
- **Present Bit(P)**: Indicating whether this page is in physical memory or on disk(swapped out)
- **Dirty Bit(D)**: Indicating whether the page has been modified since it was brought into memory
- **Reference Bit(Accessed Bit -A)**: Indicating that a page has been accessed

What is the issue with Paging?



What is the issue with Paging?

- **Size – Page tables and entries take up a lot of space in memory**




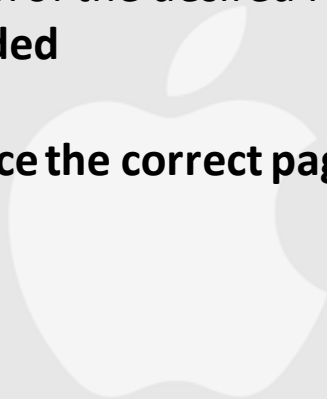

What is the issue with Paging?

- **Size** – Page tables and entries take up a lot of space in memory
- **Slow** – 2 references per memory references



Paging is Slow?

- We need to first reference the correct page table (**1 reference**)
 - To find a location of the desired PTE, the **starting location** of the page table is **needed**
- Then we can reference the correct page table entry (**1 reference**)



Paging is Slow?

- **We need to first reference the correct page table (1 reference)**
 - To find a location of the desired PTE, the **starting location** of the page table is **needed**
- **Then we can reference the correct page table entry (1 reference)**
- **Therefore, for every memory reference, paging requires the OS to perform one extra memory reference**

Accessing Memory with Paging

```
1  // Extract the VPN from the virtual address
2  VPN = (VirtualAddress & VPN_MASK) >> SHIFT

3  // Form the address of the page-table entry (PTE)
4  PTEAddr = PTBR + (VPN * sizeof(PTE))

5  // Check if process can access the page
6  if (PTE.Valid == False)
7      RaiseException(SEGMENTATION_FAULT)

8  else if (CanAccess(PTE.ProtectBits) == False)
9      RaiseException(PROTECTION_FAULT)

10 else
11     // Access is OK: form physical address and fetch it
12     offset = VirtualAddress & OFFSET_MASK
13     PhysAddr = (PTE.PFN << PFN_SHIFT) | offset
14     Register = AccessMemory(PhysAddr)
```

Accessing Memory with Paging

Example : A Simple Memory Access

```
int array[1000];  
...  
for(i = 0; i < 1000; i++)  
    array[i] = 0;
```


Compile and Execute

```
prompt> gcc -o array array.c -Wall -o  
prompt> ./array
```



Resulting Assembly code

```
0x1024 : movl $0x0, (%edi,%eax,4)  
0x1028 : incl %eax  
0x102c : cmpl $0x03e8,%eax  
0x1030 : jne 0x1024
```


A Virtual(And Physical) Memory Trace




A Virtual(And Physical) Memory Trace




A Virtual(And Physical) Memory Trace



A Virtual(And Physical) Memory Trace



A Virtual(And Physical) Memory Trace



Paging is Slow?

- We need to first reference the correct page table (**1 reference**)
 - To find a location of the desired PTE, the **starting location** of the page table is **needed**
- Then we can reference the correct page table entry (**1 reference**)
- Therefore, for every memory reference, paging requires the OS to perform one extra memory reference
- **How can we improve this access speed**


Translation Lookaside Buffers

- Part of the chip's memory-management unit(MMU).




Translation Lookaside Buffers

- Part of the chip's memory-management unit(MMU).

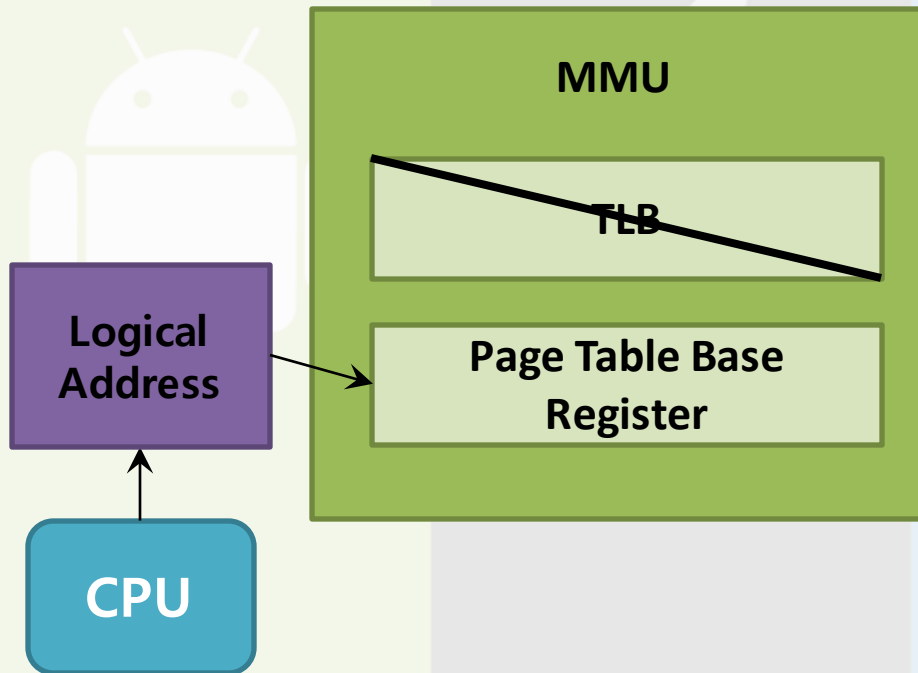


Translation Lookaside Buffers



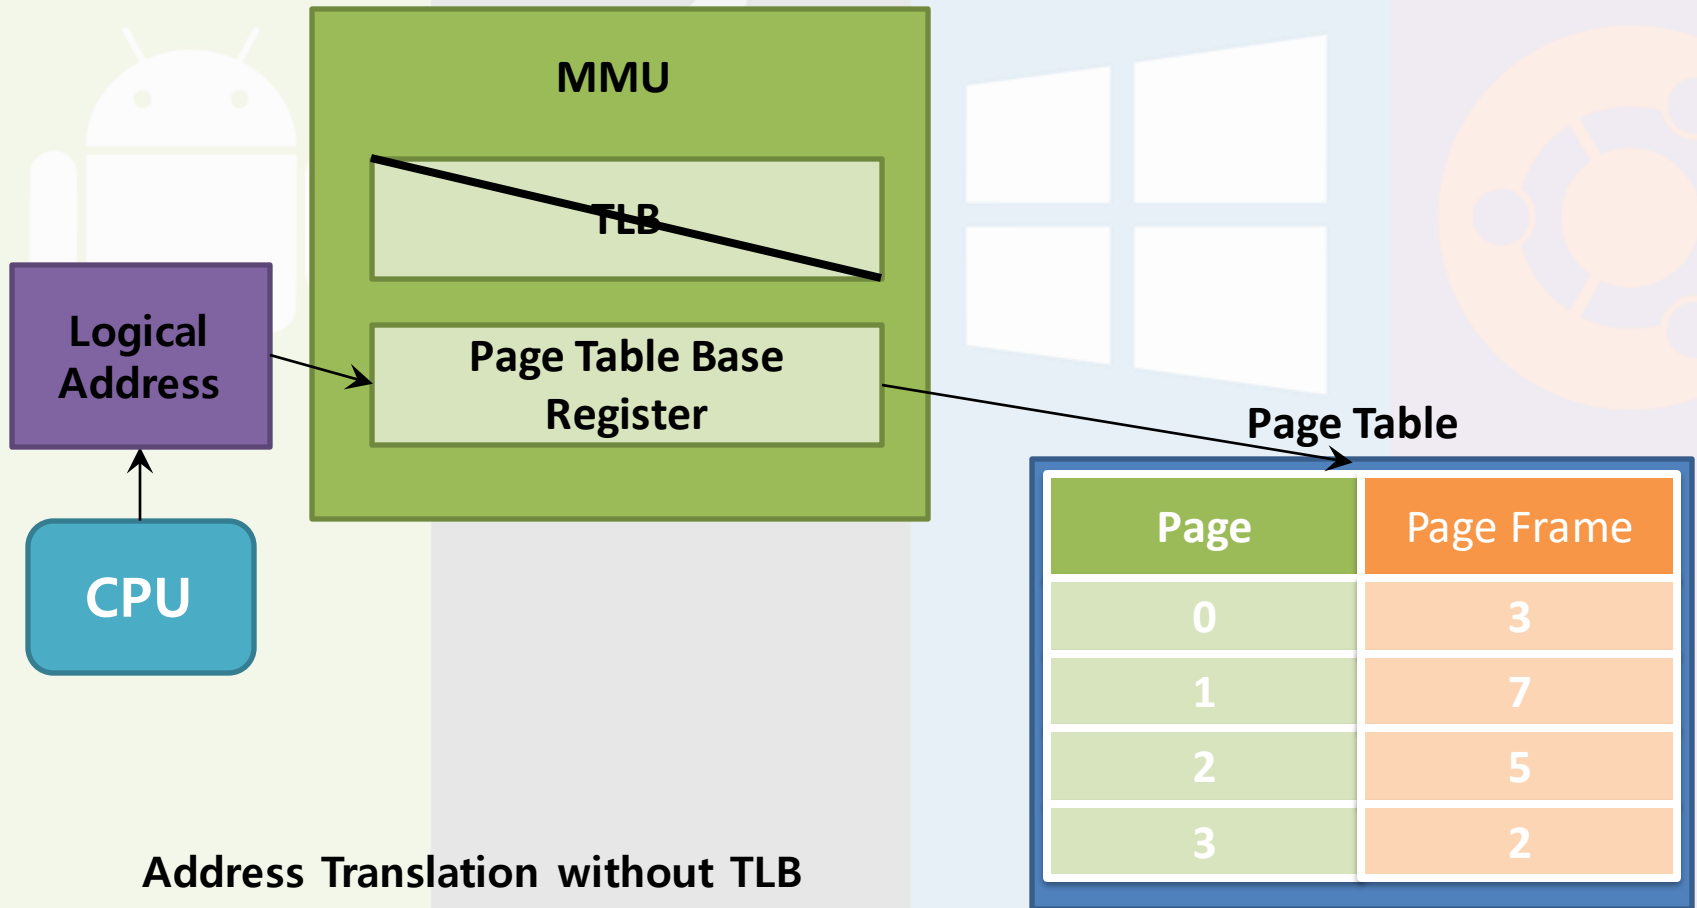
Address Translation without TLB

Translation Lookaside Buffers

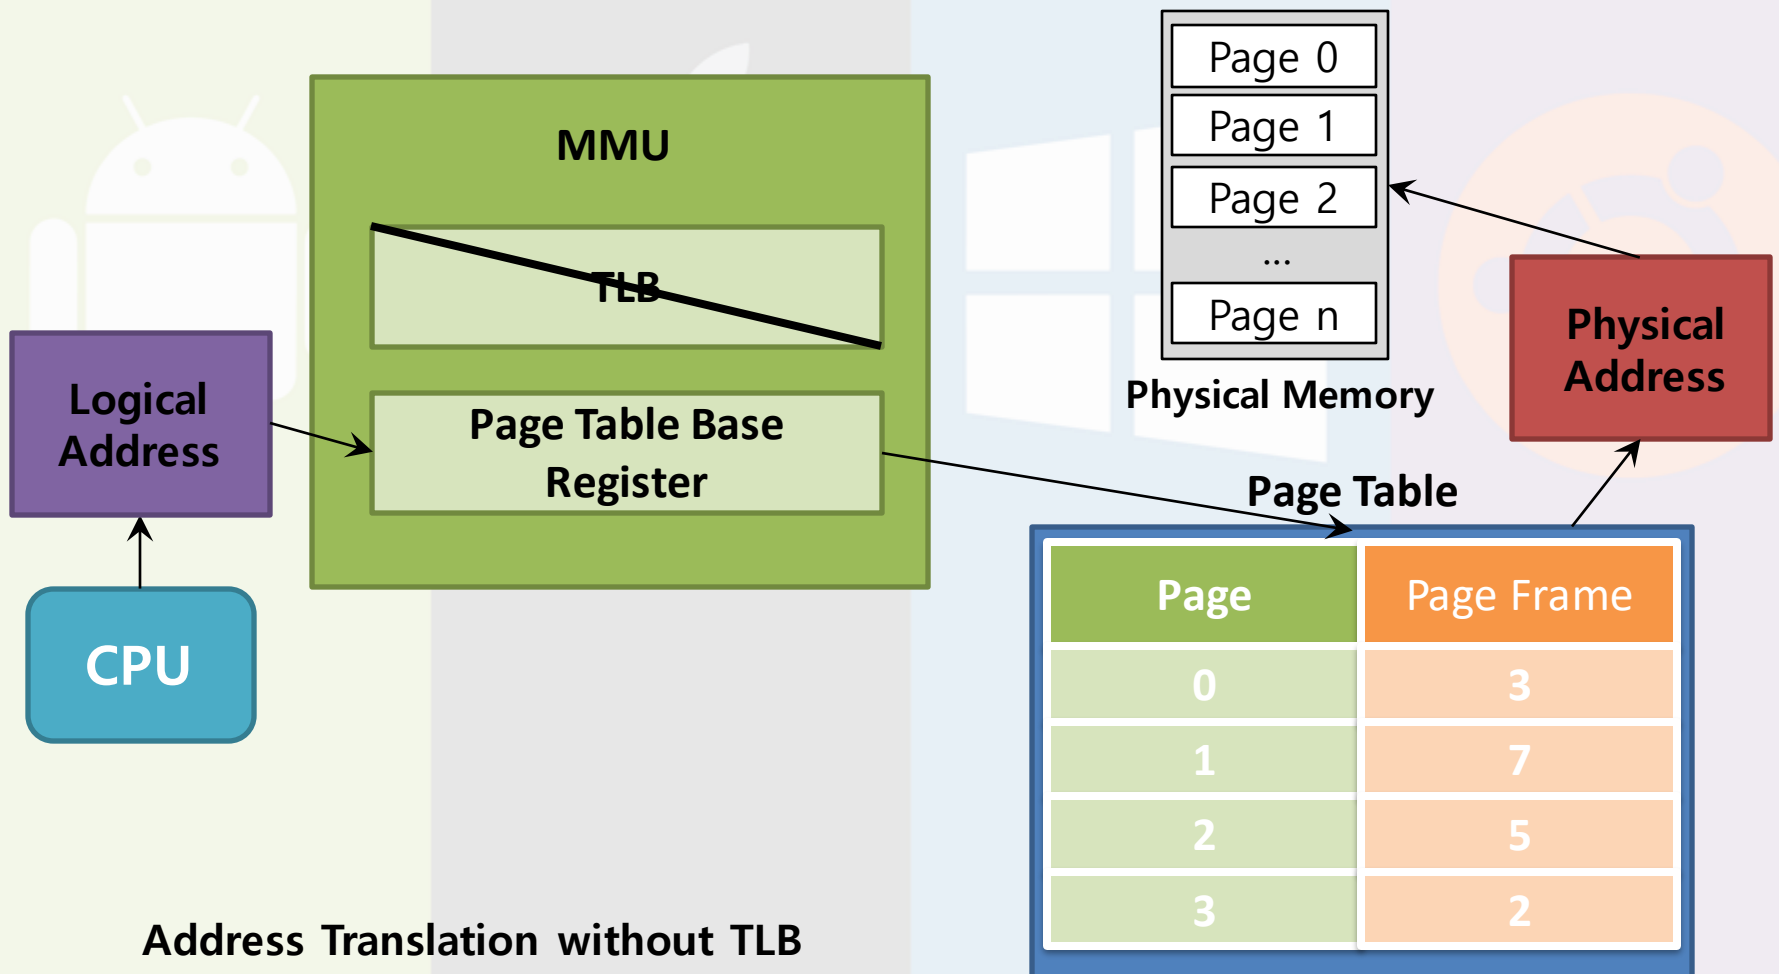


Address Translation without TLB

Translation Lookaside Buffers




Translation Lookaside Buffers



Translation Lookaside Buffers

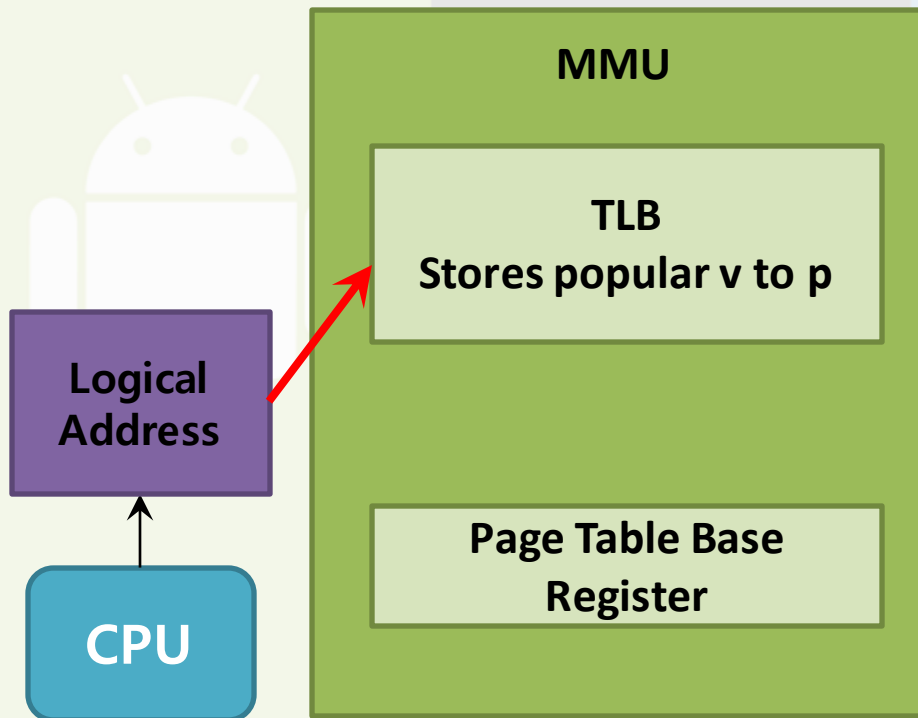
- The **TLB** is a **hardware cache** of popular virtual-to-physical address translation



Address Translation with TLB

Translation Lookaside Buffers

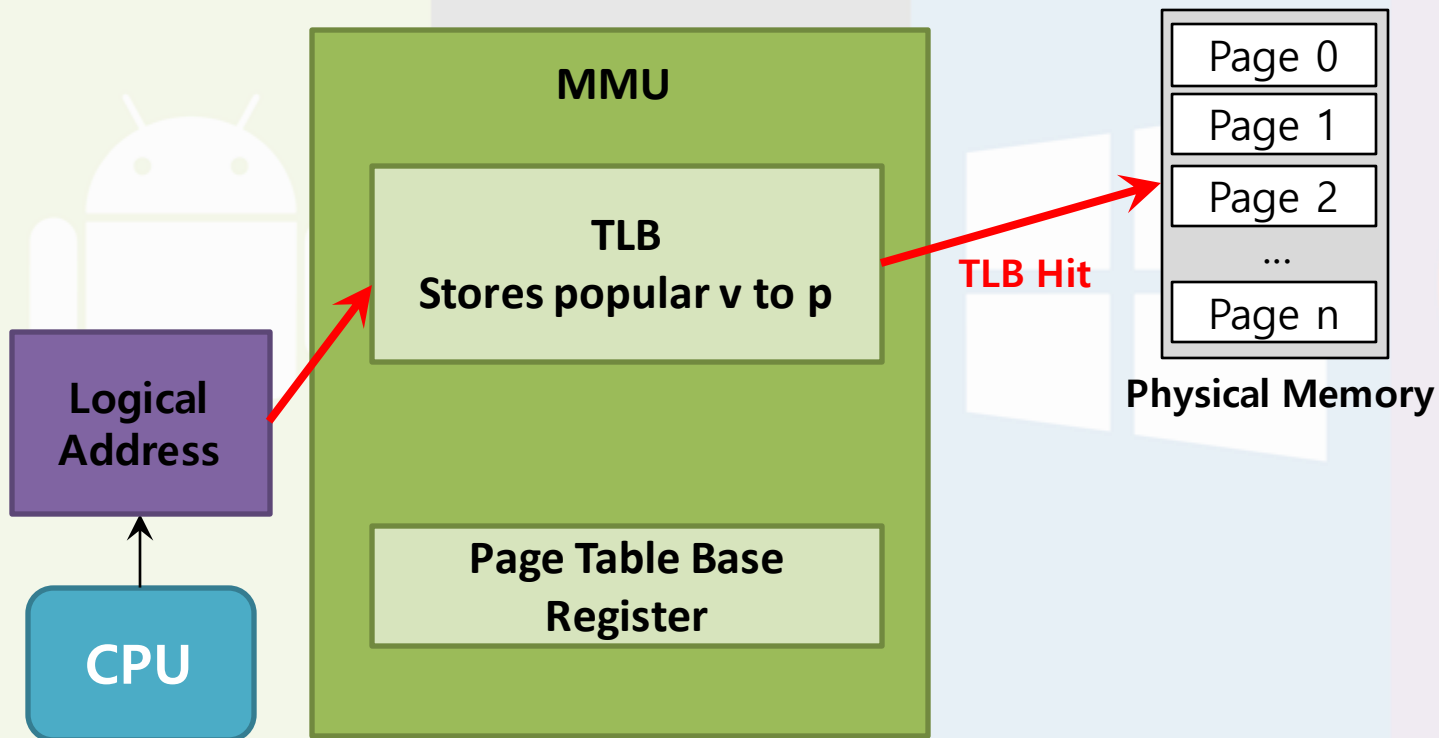
- The **TLB** is a **hardware cache** of popular virtual-to-physical address translation



Address Translation with TLB

Translation Lookaside Buffers

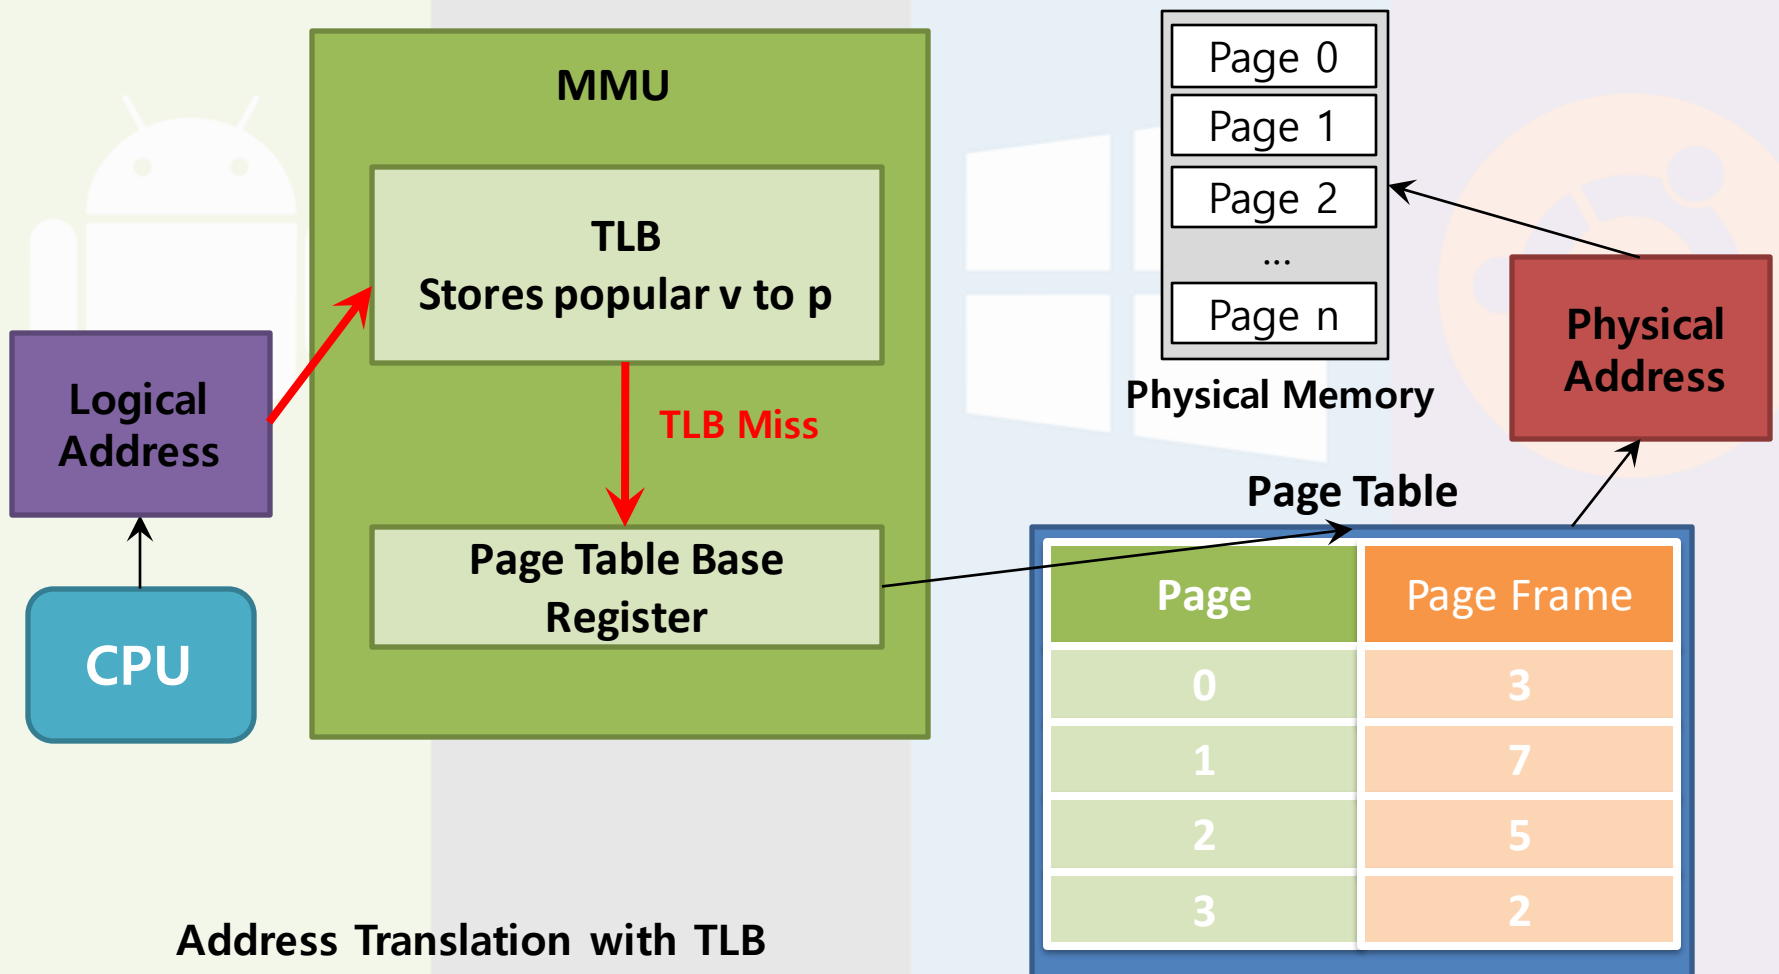
- The **TLB** is a **hardware cache** of popular virtual-to-physical address translation



Address Translation with TLB

Translation Lookaside Buffers

- The **TLB** is a **hardware cache** of popular virtual-to-physical address translation



TLB Algoritihm

```
1: VPN = (VirtualAddress & VPN_MASK ) >> SHIFT
2: (Success , TlbEntry) = TLB_Lookup(VPN)
3:   if(Success == Ture){ // TLB Hit
4:       if(CanAccess(TlbEntry.ProtectBit) == True ){
5:           offset = VirtualAddress & OFFSET_MASK
6:           PhysAddr = (TlbEntry.PFN << SHIFT) | Offset
7:           AccessMemory( PhysAddr )
8:       }else RaiseException(PROTECTION_ERROR)
```

- (1) extract the virtual page number(VPN).
- (2) check if the TLB holds the translation for this VPN.
- (5-8) extract the page frame number from the relevant TLB entry, and form the desired physical address and access memory.

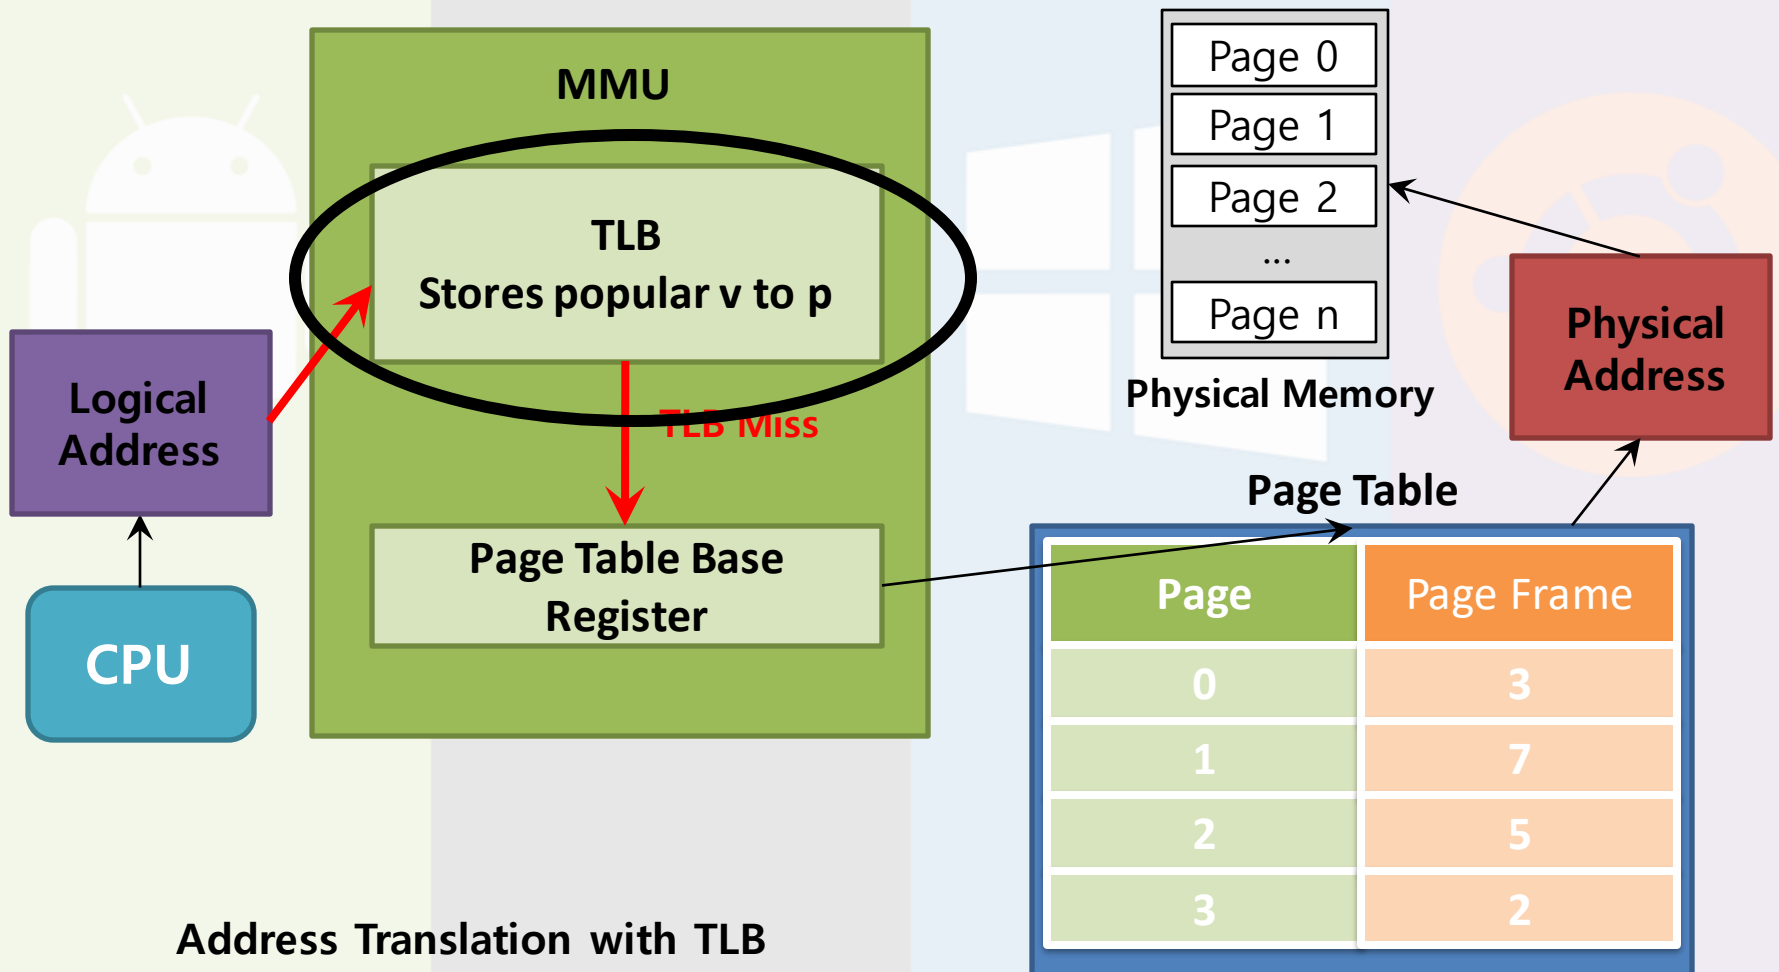
TLB Algorithm

```
11:    }else{ //TLB Miss
12:        PTEAddr = PTBR + (VPN * sizeof(PTE))
13:        PTE = AccessMemory(PTEAddr)
14:        if(PTE.VALID == FALSE){
15:            RaiseException(SEGMENTATION_FAULT)
16:        else if (CanAccess(PTE.ProtectBits) == False){
17:            RaiseException(PROTECTION_FAULT)
18:        }else{
19:            TLB_Insert( VPN , PTE.PFN , PTE.ProtectBits)
20:            RetryInstruction()
21:        }
22: }
```

- (12-13) The hardware accesses the page table to find the translation.
- (19) updates the TLB with the translation.

A TLB Entry

- The **TLB** is a **hardware cache** of popular virtual-to-physical address translation



A TLB Entry



- **TLB is managed by Full Associative method.**
 - A typical TLB might have 32, 64, or 128 entries.
 - Hardware search the entire TLB in parallel to find the desired translation.
 - Other bits: valid bits , protection bits, address-space identifier, dirty bit



Typical TLB entry look like this

Example : Accessing an Array

```
0:      int sum = 0 ;  
1:      for( i=0; i<10; i++){  
2:          sum+=a[i] ;  
3:      }
```



Example : Accessing an Array

| | OFFSET | | | | |
|----------|--------|------|------|------|----|
| | 00 | 04 | 08 | 12 | 16 |
| VPN = 00 | | | | | |
| VPN = 01 | | | | | |
| VPN = 03 | | | | | |
| VPN = 04 | | | | | |
| VPN = 05 | | | | | |
| VPN = 06 | | a[0] | a[1] | a[2] | |
| VPN = 07 | a[3] | a[4] | a[5] | a[6] | |
| VPN = 08 | a[7] | a[8] | a[9] | | |
| VPN = 09 | | | | | |
| VPN = 10 | | | | | |
| VPN = 11 | | | | | |
| VPN = 12 | | | | | |
| VPN = 13 | | | | | |
| VPN = 14 | | | | | |
| VPN = 15 | | | | | |

```
0:      int sum = 0 ;
1:      for( i=0; i<10; i++){
2:                  sum+=a[i] ;
3:      }
```

- When the first array element (a[0]) is accessed, the CPU will see a load to virtual address 100

Example : Accessing an Array

| | OFFSET | | | | |
|----------|--------|------|------|------|----|
| | 00 | 04 | 08 | 12 | 16 |
| VPN = 00 | | | | | |
| VPN = 01 | | | | | |
| VPN = 03 | | | | | |
| VPN = 04 | | | | | |
| VPN = 05 | | | | | |
| VPN = 06 | | a[0] | a[1] | a[2] | |
| VPN = 07 | a[3] | a[4] | a[5] | a[6] | |
| VPN = 08 | a[7] | a[8] | a[9] | | |
| VPN = 09 | | | | | |
| VPN = 10 | | | | | |
| VPN = 11 | | | | | |
| VPN = 12 | | | | | |
| VPN = 13 | | | | | |
| VPN = 14 | | | | | |
| VPN = 15 | | | | | |

```
0:      int sum = 0 ;
1:      for( i=0; i<10; i++){
2:                  sum+=a[i] ;
3:      }
```

- When the first array element (a[0]) is accessed, the CPU will see a load to virtual address 100
- The hardware extracts the VPN from this (VPN=06), and uses that to check the TLB for a valid translation.

Example : Accessing an Array

| | OFFSET | | | | |
|----------|--------|------|------|------|----|
| | 00 | 04 | 08 | 12 | 16 |
| VPN = 00 | | | | | |
| VPN = 01 | | | | | |
| VPN = 03 | | | | | |
| VPN = 04 | | | | | |
| VPN = 05 | | | | | |
| VPN = 06 | | a[0] | a[1] | a[2] | |
| VPN = 07 | a[3] | a[4] | a[5] | a[6] | |
| VPN = 08 | a[7] | a[8] | a[9] | | |
| VPN = 09 | | | | | |
| VPN = 10 | | | | | |
| VPN = 11 | | | | | |
| VPN = 12 | | | | | |
| VPN = 13 | | | | | |
| VPN = 14 | | | | | |
| VPN = 15 | | | | | |

| TLB | | |
|-----|-----|------------|
| VPN | PFN | other bits |
| 06 | 03 | - |

- When the first array element (a[0]) is accessed, the CPU will see a load to virtual address 100
- The hardware extracts the VPN from this (VPN=06), and uses that to check the TLB for a valid translation.
- Assuming this is the first time the program accesses the array, the result will be a **TLB miss**.
- Add new TLB Entry

Example : Accessing an Array

| | OFFSET | | | | |
|----------|--------|------|------|------|----|
| | 00 | 04 | 08 | 12 | 16 |
| VPN = 00 | | | | | |
| VPN = 01 | | | | | |
| VPN = 03 | | | | | |
| VPN = 04 | | | | | |
| VPN = 05 | | | | | |
| VPN = 06 | | a[0] | a[1] | a[2] | |
| VPN = 07 | a[3] | a[4] | a[5] | a[6] | |
| VPN = 08 | a[7] | a[8] | a[9] | | |
| VPN = 09 | | | | | |
| VPN = 10 | | | | | |
| VPN = 11 | | | | | |
| VPN = 12 | | | | | |
| VPN = 13 | | | | | |
| VPN = 14 | | | | | |
| VPN = 15 | | | | | |

| TLB | | |
|-----|-----|------------|
| VPN | PFN | other bits |
| 06 | 03 | - |

- The next access is to a[1]
- The hardware extracts the VPN from this (VPN=06), and uses that to check the TLB for a valid translation.

Example : Accessing an Array

| | OFFSET | | | | |
|----------|--------|------|------|------|----|
| | 00 | 04 | 08 | 12 | 16 |
| VPN = 00 | | | | | |
| VPN = 01 | | | | | |
| VPN = 03 | | | | | |
| VPN = 04 | | | | | |
| VPN = 05 | | | | | |
| VPN = 06 | | a[0] | a[1] | a[2] | |
| VPN = 07 | a[3] | a[4] | a[5] | a[6] | |
| VPN = 08 | a[7] | a[8] | a[9] | | |
| VPN = 09 | | | | | |
| VPN = 10 | | | | | |
| VPN = 11 | | | | | |
| VPN = 12 | | | | | |
| VPN = 13 | | | | | |
| VPN = 14 | | | | | |
| VPN = 15 | | | | | |

| TLB | | |
|-----|-----|------------|
| VPN | PFN | other bits |
| 06 | 03 | - |

- The next access is to a[1]
- The hardware extracts the VPN from this (VPN=06), and uses that to check the TLB for a valid translation.
- Some good news here: a **TLB hit!** Because the second element of the array is packed next to the first, it lives on the same page which had already been loaded into the TLB

Example : Accessing an Array

| | OFFSET | | | | |
|----------|--------|------|------|------|----|
| | 00 | 04 | 08 | 12 | 16 |
| VPN = 00 | | | | | |
| VPN = 01 | | | | | |
| VPN = 03 | | | | | |
| VPN = 04 | | | | | |
| VPN = 05 | | | | | |
| VPN = 06 | | a[0] | a[1] | a[2] | |
| VPN = 07 | a[3] | a[4] | a[5] | a[6] | |
| VPN = 08 | a[7] | a[8] | a[9] | | |
| VPN = 09 | | | | | |
| VPN = 10 | | | | | |
| VPN = 11 | | | | | |
| VPN = 12 | | | | | |
| VPN = 13 | | | | | |
| VPN = 14 | | | | | |
| VPN = 15 | | | | | |

| TLB | | |
|-----|-----|------------|
| VPN | PFN | other bits |
| 06 | 03 | - |

- The next access is to a[2]
- (VPN=06)
- **TLB hit!**

Example : Accessing an Array

| | OFFSET | | | | |
|----------|--------|------|------|------|----|
| | 00 | 04 | 08 | 12 | 16 |
| VPN = 00 | | | | | |
| VPN = 01 | | | | | |
| VPN = 03 | | | | | |
| VPN = 04 | | | | | |
| VPN = 05 | | | | | |
| VPN = 06 | | a[0] | a[1] | a[2] | |
| VPN = 07 | a[3] | a[4] | a[5] | a[6] | |
| VPN = 08 | a[7] | a[8] | a[9] | | |
| VPN = 09 | | | | | |
| VPN = 10 | | | | | |
| VPN = 11 | | | | | |
| VPN = 12 | | | | | |
| VPN = 13 | | | | | |
| VPN = 14 | | | | | |
| VPN = 15 | | | | | |

| TLB | | |
|-----|-----|------------|
| VPN | PFN | other bits |
| 06 | 03 | - |
| 07 | 08 | - |

- The next access is to a[3]
- (VPN=07)
- **TLB miss!**
- Load translation into TLB

Example : Accessing an Array

| | OFFSET | | | | |
|----------|--------|------|------|------|----|
| | 00 | 04 | 08 | 12 | 16 |
| VPN = 00 | | | | | |
| VPN = 01 | | | | | |
| VPN = 03 | | | | | |
| VPN = 04 | | | | | |
| VPN = 05 | | | | | |
| VPN = 06 | | a[0] | a[1] | a[2] | |
| VPN = 07 | a[3] | a[4] | a[5] | a[6] | |
| VPN = 08 | a[7] | a[8] | a[9] | | |
| VPN = 09 | | | | | |
| VPN = 10 | | | | | |
| VPN = 11 | | | | | |
| VPN = 12 | | | | | |
| VPN = 13 | | | | | |
| VPN = 14 | | | | | |
| VPN = 15 | | | | | |

| TLB | | |
|-----|-----|------------|
| VPN | PFN | other bits |
| 06 | 03 | - |
| 07 | 08 | - |

- The next access is to a[4]
- (VPN=07)
- **TLB hit!**

Example : Accessing an Array

| | OFFSET | | | | |
|----------|--------|------|------|------|----|
| | 00 | 04 | 08 | 12 | 16 |
| VPN = 00 | | | | | |
| VPN = 01 | | | | | |
| VPN = 03 | | | | | |
| VPN = 04 | | | | | |
| VPN = 05 | | | | | |
| VPN = 06 | | a[0] | a[1] | a[2] | |
| VPN = 07 | a[3] | a[4] | a[5] | a[6] | |
| VPN = 08 | a[7] | a[8] | a[9] | | |
| VPN = 09 | | | | | |
| VPN = 10 | | | | | |
| VPN = 11 | | | | | |
| VPN = 12 | | | | | |
| VPN = 13 | | | | | |
| VPN = 14 | | | | | |
| VPN = 15 | | | | | |

| TLB | | |
|-----|-----|------------|
| VPN | PFN | other bits |
| 06 | 03 | - |
| 07 | 08 | - |

- The next access is to a[5]
- (VPN=07)
- **TLB hit!**

Example : Accessing an Array

| | OFFSET | | | | |
|----------|--------|------|------|------|----|
| | 00 | 04 | 08 | 12 | 16 |
| VPN = 00 | | | | | |
| VPN = 01 | | | | | |
| VPN = 03 | | | | | |
| VPN = 04 | | | | | |
| VPN = 05 | | | | | |
| VPN = 06 | | a[0] | a[1] | a[2] | |
| VPN = 07 | a[3] | a[4] | a[5] | a[6] | |
| VPN = 08 | a[7] | a[8] | a[9] | | |
| VPN = 09 | | | | | |
| VPN = 10 | | | | | |
| VPN = 11 | | | | | |
| VPN = 12 | | | | | |
| VPN = 13 | | | | | |
| VPN = 14 | | | | | |
| VPN = 15 | | | | | |

| TLB | | |
|-----|-----|------------|
| VPN | PFN | other bits |
| 06 | 03 | - |
| 07 | 08 | - |

- The next access is to a[6]
- (VPN=07)
- **TLB hit!**

Example : Accessing an Array

| | OFFSET | | | | |
|----------|--------|------|------|------|----|
| | 00 | 04 | 08 | 12 | 16 |
| VPN = 00 | | | | | |
| VPN = 01 | | | | | |
| VPN = 03 | | | | | |
| VPN = 04 | | | | | |
| VPN = 05 | | | | | |
| VPN = 06 | | a[0] | a[1] | a[2] | |
| VPN = 07 | a[3] | a[4] | a[5] | a[6] | |
| VPN = 08 | a[7] | a[8] | a[9] | | |
| VPN = 09 | | | | | |
| VPN = 10 | | | | | |
| VPN = 11 | | | | | |
| VPN = 12 | | | | | |
| VPN = 13 | | | | | |
| VPN = 14 | | | | | |
| VPN = 15 | | | | | |

| TLB | | |
|-----|-----|------------|
| VPN | PFN | other bits |
| 06 | 03 | - |
| 07 | 08 | - |
| 08 | 02 | - |

- The next access is to a[7]
- (VPN=08)
- **TLB miss!**
- Load translation into TLB

Example : Accessing an Array

| | OFFSET | | | | |
|----------|--------|------|------|------|----|
| | 00 | 04 | 08 | 12 | 16 |
| VPN = 00 | | | | | |
| VPN = 01 | | | | | |
| VPN = 03 | | | | | |
| VPN = 04 | | | | | |
| VPN = 05 | | | | | |
| VPN = 06 | | a[0] | a[1] | a[2] | |
| VPN = 07 | a[3] | a[4] | a[5] | a[6] | |
| VPN = 08 | a[7] | a[8] | a[9] | | |
| VPN = 09 | | | | | |
| VPN = 10 | | | | | |
| VPN = 11 | | | | | |
| VPN = 12 | | | | | |
| VPN = 13 | | | | | |
| VPN = 14 | | | | | |
| VPN = 15 | | | | | |

| TLB | | |
|-----|-----|------------|
| VPN | PFN | other bits |
| 06 | 03 | - |
| 07 | 08 | - |
| 08 | 02 | - |

- The next access is to a[8]
- (VPN=08)
- **TLB hit!**

Example : Accessing an Array

| | OFFSET | | | | |
|----------|--------|------|------|------|----|
| | 00 | 04 | 08 | 12 | 16 |
| VPN = 00 | | | | | |
| VPN = 01 | | | | | |
| VPN = 03 | | | | | |
| VPN = 04 | | | | | |
| VPN = 05 | | | | | |
| VPN = 06 | | a[0] | a[1] | a[2] | |
| VPN = 07 | a[3] | a[4] | a[5] | a[6] | |
| VPN = 08 | a[7] | a[8] | a[9] | | |
| VPN = 09 | | | | | |
| VPN = 10 | | | | | |
| VPN = 11 | | | | | |
| VPN = 12 | | | | | |
| VPN = 13 | | | | | |
| VPN = 14 | | | | | |
| VPN = 15 | | | | | |

| TLB | | |
|-----|-----|------------|
| VPN | PFN | other bits |
| 06 | 03 | - |
| 07 | 08 | - |
| 08 | 02 | - |

- The next access is to a[9]
- (VPN=08)
- **TLB hit!**

Example : Accessing an Array

| | OFFSET | | | | |
|----------|--------|------|------|------|----|
| | 00 | 04 | 08 | 12 | 16 |
| VPN = 00 | | | | | |
| VPN = 01 | | | | | |
| VPN = 03 | | | | | |
| VPN = 04 | | | | | |
| VPN = 05 | | | | | |
| VPN = 06 | | a[0] | a[1] | a[2] | |
| VPN = 07 | a[3] | a[4] | a[5] | a[6] | |
| VPN = 08 | a[7] | a[8] | a[9] | | |
| VPN = 09 | | | | | |
| VPN = 10 | | | | | |
| VPN = 11 | | | | | |
| VPN = 12 | | | | | |
| VPN = 13 | | | | | |
| VPN = 14 | | | | | |
| VPN = 15 | | | | | |

| TLB | | |
|-----|-----|------------|
| VPN | PFN | other bits |
| 06 | 03 | - |
| 07 | 08 | - |
| 08 | 02 | - |

- The next access is to a[9]
- (VPN=08)
- **TLB hit!**

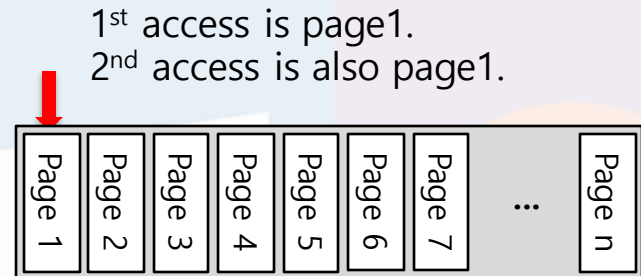
The TLB improves performance
due to **spatial locality**

3 misses and 7 hits.
Thus **TLB hit rate** is 70%.

Locality

- **Temporal Locality**

- An instruction or data item that has been recently accessed will likely be re-accessed soon in the future




Virtual Memory

Locality


- **Temporal Locality**

- An instruction or data item that has been recently accessed will likely be re-accessed soon in the future



- **Spatial Locality**

- If a program accesses memory at address x , it will likely soon access memory near x



Who handles the TLB Miss?

- Well there are 2 approaches



Who handles the TLB Miss?

- Hardware handle the TLB miss entirely on CISC



Who handles the TLB Miss?

- **Hardware handle the TLB miss entirely on CISC**
 - The hardware has to know exactly where the page tables are located in memory.
 - The hardware would “walk” the page table, find the correct page-table entry and extract the desired translation, update and retry instruction.
 - **hardware-managed TLB.**

Who handles the TLB Miss?

- **Hardware handle the TLB miss entirely on CISC**
 - The hardware has to know exactly where the page tables are located in memory.
 - The hardware would “walk” the page table, find the correct page-table entry and extract the desired translation, update and retry instruction.
 - **hardware-managed TLB.**
- **RISC have what is known as a software-managed TLB**

Who handles the TLB Miss?

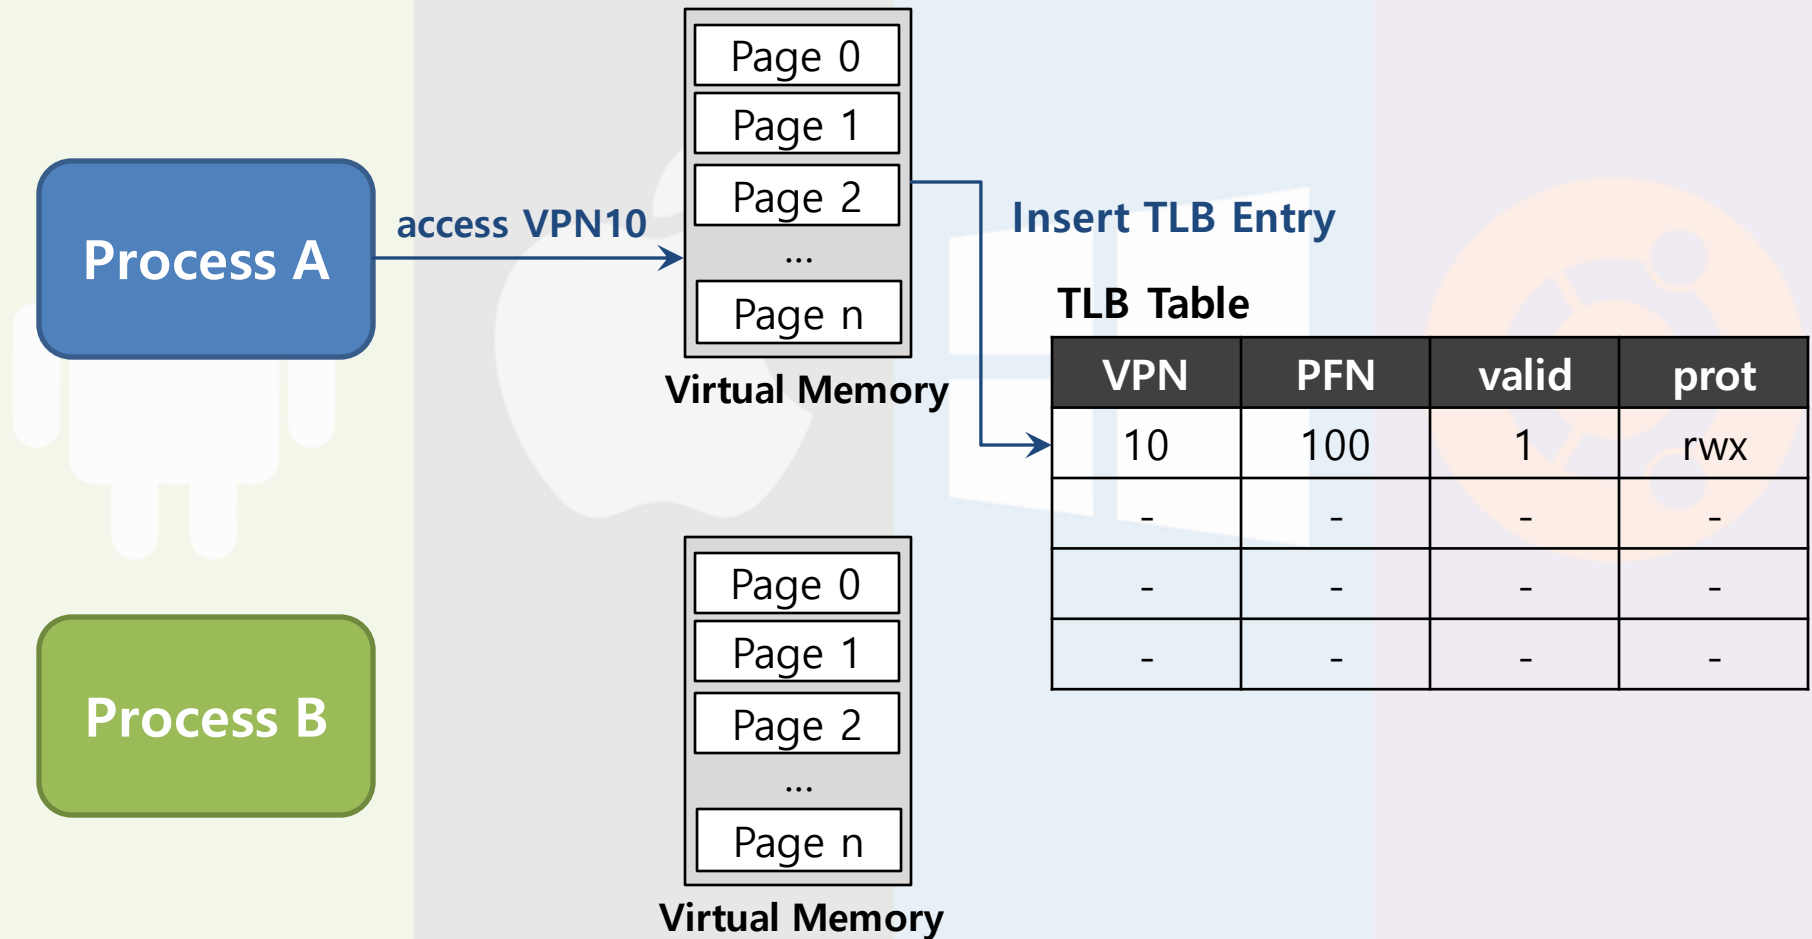
- **Hardware handle the TLB miss entirely on CISC**
 - The hardware has to know exactly where the page tables are located in memory.
 - The hardware would “walk” the page table, find the correct page-table entry and extract the desired translation, update and retry instruction.
 - **hardware-managed TLB**
 - Complex Instruction Set Computing
- **RISC have what is known as a software-managed TLB**
 - On a TLB miss, the hardware raises exception(trap handler).
 - **Trap handler is code** within the OS that is written with the express purpose of handling TLB miss.
 - Reduced instruction set computer

Ok so job Done?

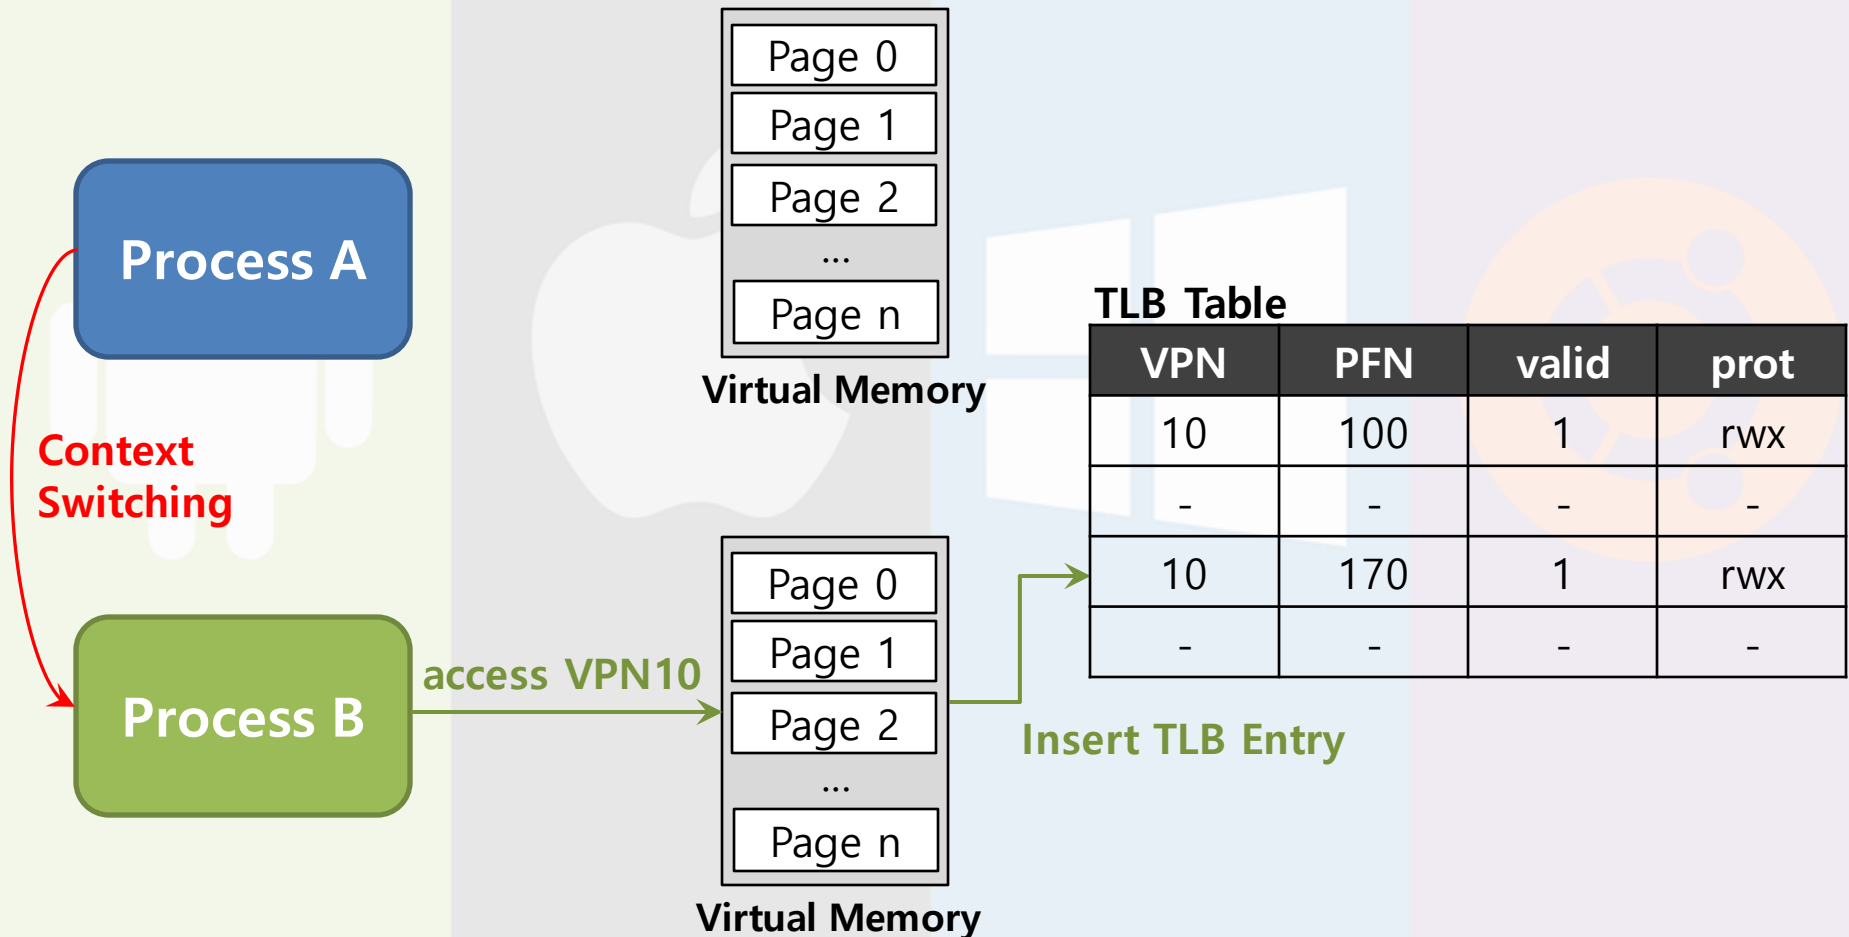
- The TLB in its current form solves all our problems



TLB Issue : Context Switching




TLB Issue : Context Switching




TLB Issue : Context Switching

Process A



Process B



TLB Table

| VPN | PFN | valid | prot |
|-----|-----|-------|------|
| 10 | 100 | 1 | rwX |
| - | - | - | - |
| 10 | 170 | 1 | rwX |
| - | - | - | - |

Can't **Distinguish** which entry is meant for which process

TLB Issue : Page Sharing

- **Two processes share a page.**
 - Process 1 is sharing physical page 101 with Process2
 - P1 maps this page into the 10th page of its address space
 - P2 maps this page to the 50th page of its address space


| VPN | PFN | valid | prot |
|-----|-----|-------|------|
| 10 | 101 | 1 | rwX |
| - | - | - | - |
| 50 | 101 | 1 | rwX |
| - | - | - | - |

Sharing of pages is **useful** as it reduces the number of physical pages in use.

TLB Issue : Context Switching - **Solution**


- Provide an address space identifier(ASID) field in the TLB

Process A



Virtual Memory

Process B



Virtual Memory

TLB Table

| VPN | PFN | valid | prot | ASID |
|-----|-----|-------|------|------|
| 10 | 100 | 1 | rwX | 1 |
| - | - | - | - | - |
| 10 | 170 | 1 | rwX | 2 |
| - | - | - | - | - |

What happens when we want to add a new entry entry

- Problem – We need to add a new entry to the TLB
 - Well where do we put it and if needed which entry do we replace



TLB Issue : Replacement Policy

- In order to add new entries a replacement policy is used
- The general idea with these policies is that you want to **minimize the miss rate** and hence **maximise the hit rate**



TLB Issue : Replacement Policy

- In order to add new entries a replacement policy is used
- The general idea with these policies is that you want to **minimize the miss rate** and hence **maximise the hit rate**




- LIFO
- FIFO
- Random
- LRU



TLB Issue : Replacement Policy - **Solution**

- LRU (LEAST RECENTLY USED)

- Evict an entry that has not recently been used.
- Take advantage of locality in the memory-reference stream



TLB Issue : Replacement Policy - **Solution**

- **LRU (LEAST RECENTLY USED)**

- **Evict an entry that has not recently been used.**
- **Take advantage of locality in the memory-reference stream**
 - Add a register to every page frame - contain the last time that the page in that frame was accessed
 - Use a "logical clock" that advance by 1 tick each time a memory reference is made.
 - Each time a page is referenced, update its register

TLB Issue : Replacement Policy - **Solution**

- **LRU (LEAST RECENTLY USED)**

- Evict an entry that has not recently been used.
- Take advantage of locality in the memory-reference stream

Page Reference Row

7 0 1 2 0 3 0 4 2 3 0 3 2 1 2 0


Page Frame:

| | | | | | | | | | | | | | | | | | | | | | | | | | | | | | | | |
|---|---|---|---|---|---|---|---|---|---|---|---|---|---|----|---|----|---|----|---|----|---|----|---|----|---|----|---|----|---|----|---|
| 1 | 7 | 1 | 7 | 1 | 7 | 4 | 2 | 4 | 2 | 4 | 2 | 4 | 2 | 8 | 4 | 8 | 4 | 8 | 4 | 11 | 0 | 11 | 0 | 11 | 0 | 14 | 1 | 14 | 1 | 14 | 1 |
| 2 | 0 | 2 | 0 | 2 | 0 | 5 | 0 | 5 | 0 | 7 | 0 | 7 | 0 | 10 | 3 | 10 | 3 | 12 | 3 | 12 | 3 | 12 | 3 | 12 | 3 | 16 | 0 | 16 | 0 | 16 | 0 |
| 3 | 1 | 3 | 1 | 3 | 1 | 6 | 3 | 6 | 3 | 6 | 3 | 6 | 3 | 9 | 2 | 9 | 2 | 9 | 2 | 9 | 2 | 9 | 2 | 13 | 2 | 13 | 2 | 15 | 2 | 15 | 2 |

Total 11 TLB miss

A Real TLB Entry

All 64 bits of this TLB entry(example of MIPS R4000)



| Flag | Content |
|------------------|--|
| 19-bit VPN | The rest reserved for the kernel. |
| 24-bit PFN | Systems can support with up to 64GB of main memory($2^{24} * 4KB$ pages) |
| Global bit(G) | Used for pages that are globally-shared among processes. |
| ASID | OS can use to distinguish between address spaces. |
| Coherence bit(C) | determine how a page is cached by the hardware. |
| Dirty bit(D) | marking when the page has been written. |
| Valid bit(V) | tells the hardware if there is a valid translation present in the entry. |